



ĐẠI HỌC BÁCH KHOA HÀ NỘI

Course: Big Data Storage and Processing

Real-Time Stock Sentiment Analysis Pipeline

Group 14

Le Ngoc Binh	20214878
Hoang Trung Khai	20225502
Nguyen Viet Anh	20225434
Trinh Duy Phong	20220065
Luu Hoang Phan	20225516



Problem Statement


Objectives

- Build a complete **end-to-end** system capable of **collecting, processing, and performing real-time sentiment analysis** on **social media data** (Twitter/X).

Key Highlights

- Build a real-time streaming data pipeline
- Microservices-based architecture for modularity and flexibility.
- Process data using **Kafka + Spark Streaming**
- Fully containerized deployment using **Docker**.
- **Kubernetes (Minikube)** orchestration for scalability and reliability.
- Infrastructure managed as Code (IaC) with **Terraform** for consistent and repeatable deployments.
- Visualize insights using **Grafana dashboards**

Overview



I. Data Ingestion

This step collects and generates data into the system

Twitter Producer

- Stimulates live stock-related tweets
- Reads historical dataset (Hugging Face)
- Overwrites timestamps with current UTC time
- Sends JSON messages to Kafka topic tweets

1

Stock Producer

- Simulates minute-level OHLCV stock prices
- Reads static CSV data
- Generates real-time trading minutes
- Sends data to Kafka topic stock-prices

2

I. Data Ingestion

This step collects and generates data into the system

Twitter Producer

- Stimulates live stock-related tweets
- Reads historical dataset (Hugging Face)
- Overwrites timestamps with current UTC time
- Sends JSON messages to Kafka topic tweets

1

Stock Producer

- Simulates minute-level OHLCV stock prices
- Reads static CSV data
- Generates real-time trading minutes
- Sends data to Kafka topic stock-prices

2

Kafka Cluster

- Acts as a **buffer and decoupling layer** on **Kubernetes with Zookeeper**
- Ensures scalability and fault tolerance
- Topics: *tweets, stock-prices*

3

Data Flow

Stage 1 - Generation

Producers generate simulated real-time data


Tweet JSON

```
{  
  "created_at": "2025-12-21T14:30:00.123456",  
  "text": "Huge breakout for $AAPL today! Testing new  
highs."  
}
```

Stock Price JSON

```
{  
  "symbol": "AAPL",  
  "trade_timestamp": "2025-12-21T14:30:00",  
  "open_price": 150.5,  
  "close_price": 151.2,  
  "high_price": 151.5,  
  "low_price": 150.0,  
  "volume": 551923  
}
```

Data Flow




Stage 1 - Generation

Producers generate simulated real-time data

Stage 2 - Buffering

Incoming messages are buffered in **Kafka Topics** (*tweets, stock-prices*) under UTF-8 JSON.
The data structure is exactly the same as in previous step.

Data Flow



Stage 1 - Generation

Producers generate simulated real-time data

Stage 2 - Buffering

Incoming messages are buffered in **Kafka Topics** (*tweets*, *stock-prices*) under UTF-8 JSON.
The data structure is exactly the same as in previous step.

Stage 3 - Processing

Spark consumes and processes streams

Spark reads JSON data, converts it into a DataFrame, and doing ETL & Analytics


Step 1: Raw Parsing (Tweet)

- Extract symbol: \$AAPL → AAPL
- Convert created_at → api_timestamp

Step 2: Compute *sentiment_score*

Step 3: Aggregation: Data is aggregated by time window and topic.

Data Flow



Stage 1 - Generation

Producers generate simulated real-time data

Stage 2 - Buffering

Incoming messages are buffered in **Kafka Topics** (*tweets*, *stock-prices*) under UTF-8 JSON.
The data structure is exactly the same as in previous step.

Stage 3 - Processing


Spark consumes and processes streams

Spark reads JSON data, converts it into a DataFrame, and doing ETL & Analytics

Stage 4 - Storage

Store the data in MinIO (Data Lake) and Cassandra (Serving Layer).

Data Flow



Stage 1 - Generation

Producers generate simulated real-time data

Stage 2 - Buffering

Incoming messages are buffered in **Kafka Topics** (*tweets*, *stock-prices*) under UTF-8 JSON.
The data structure is exactly the same as in previous step.

Stage 3 - Processing

Spark consumes and processes streams
Spark reads JSON data, converts it into a DataFrame, and doing ETL & Analytics

Stage 4 - Storage

Store the data in MinIO (Data Lake) and Cassandra (Serving Layer).

Stage 5 - Visualization

Grafana queries Cassandra to plot charts and visualize the data

II. Speed Layer

- Engine: **Apache Spark Structured Streaming 3.4.0 (Python/PySpark)**.
- Act as the backbone of the system, where raw data is transformed into information.

Input:

tweets	stock-prices
<p>Structure:</p> <ul style="list-style-type: none">• created_at: StringType• text: StringType	<p>Structure:</p> <ul style="list-style-type: none">• symbol: StringType• trade_timestamp: StringType• open_price: DoubleType• high_price: DoubleType• low_price: DoubleType• close_price: DoubleType• volume: LongType

II. Speed Layer

Spark performs a series of transformations on the data stream (DStream/DataFrame):

01

Read & Parse input

Read data from Kafka and decode JSON into columns.

02

Enrichment & Sentiment Analysis

Use RegEx to extract stock symbols
Calculate Sentiment
Create a unique ID for each record

03

Window Aggregation

Aggregate data by time windows (5-minute window, 1-minute sliding) and by topic.

Metrics:

- Average sentiment score
- Bullish tweet count
- Bearish tweet count
- Neutral tweet count

II. Speed Layer

Streaming Writers (Output)

- **Writer 1 (Bronze - MinIO):** Writes all `enriched_tweets` to MinIO in Parquet format, partitioned by `topic`.
- **Writer 2 (Speed - Cassandra Raw):** Writes each individual tweet record to the `twitter.tweets` table.
- **Writer 3 (Speed - Cassandra Market):** Writes stock price data to the `twitter.market_data` table.
- **Writer 5 (Gold - Cassandra Analytics):** Writes aggregated results (`windowed_sentiment`) to the `twitter.topic_sentiment_avg` table.

III. Batch Processing Layer

Overview

- The analytical engine of our Lambda Architecture, responsible for deep, historical analysis of "Big Data."
- Corrects potential errors or gaps from the real-time stream layer.

Input & Technology

- Engine : Apache Spark 3.4 (PySpark)
- Input Source : MinIO (S3 Compatible Data Lake)
- Data Format : Parquet (A highly efficient columnar storage format)

III. Batch Processing Layer

This layer performs three heavy-duty analyses:

01

Market Heatmap

Groups millions of data points by the hour and pivots the data to create a sentiment matrix with time slots as rows and stock symbols as columns.

02

Volatility & Risk Profiling

Calculates Standard Deviation to measure volatility and uses Median and 95th Percentile to identify the true central tendency, ignoring outliers.

03

Long-Term Trend Identification

Applies a Moving Average over a sliding window to smooth out noisy real-time data and reveal the underlying market trend

III. Batch Processing Layer

Desired Output (The "Gold" Layer)


The processed, high-value data is stored in a "Gold" layer in our data lake, ready for BI tools and dashboards.

Performance Optimizations

- **In-Memory Caching :** The source dataset is loaded into RAM once and reused for all three transformations, avoiding multiple slow disk reads.
- **Partitioning :** Results are written back to MinIO partitioned by topic (e.g., AAPL , GOOGL). This allows for highly efficient queries, as analysts can read data for a single stock without scanning the entire dataset.

III. Storage Layer

- The system uses a **Hybrid Storage Architecture**
- Designed to support:
 - Fast real-time queries (Dashboards)
 - Long-term storage (Analytics & ML)
- Core technologies:
 - **Apache Cassandra** – Speed / Serving Layer
 - **MinIO** – Data Lake (Bronze Layer)



cassandra



III. Storage Layer

3.1. Apache Cassandra (Speed Layer / OLTP Database)

- **Role:**
 - High-performance NoSQL database
 - Serves **real-time queries** from Grafana dashboards
- **Key characteristics:**
 - Low latency
 - Horizontally scalable
- **Keyspace:**
 - `twitter`
 - Replication Factor: 1 (development environment)

III. Storage Layer

3.1. Apache Cassandra (Speed Layer / OLTP Database)

Cassandra Table:

- ***tweets***: Store individual tweets after enrichment, used for detailed inspection
- ***market_data***: Store market price data, enables correlation between sentiment and price
- ***topic_sentiment_avg***: Store windowed sentiment analytics. This is data source for Grafana charts

III. Storage Layer

Apache Cassandra Table Examples

tweets

api_timestamp	sentiment_score	text
2025-12-21 07:22:05.727000+0000	0.4404	\$AAPL truly the scumbaggiest stock that ever existed

market_data

symbol	trade_timestamp	close_price	volume
AAPL	2025-12-22 13:54:13.000000+0000	68.39574	551923

topic_sentiment_avg

topic	ingest_timestamp	sentiment_score_avg	bullish_count	bearish_count	neutral_count
AAPL	2025-12-21 07:18:00.000000+0000	0.06764	1	0	4

III. Storage Layer

2. MinIO (Data Lake)

- S3-compatible object storage
- **Stores:** Raw & semi-processed tweet data
- **Format:** Parquet
- **Used for:**
 - Batch analytics
 - Data science & ML workloads
- Data is automatically partitions by topic

III. Storage Layer

2. MinIO (Data Lake)

Example Structure

```
/data/twitter-bronze/tweets_v2/
    ├── topic=AAPL/
    │   ├── part-00000-....c000.snappy.parquet (created 2025-12-21 07:xx:xx)
    │   ├── part-00001-....c000.snappy.parquet (created 2025-12-21 07:yy:yy)
    ├── topic=GOOGL/
    │   ├── part-00000-....c000.snappy.parquet
    └── _spark_metadata/
```

IV. Serving Layer

- The end-user interface for monitoring the market.
- Built with **Grafana** on top of **Apache Cassandra**

Grafana Setup

- URL: *http://localhost:3000* (via Port Forwarding)
- Datasource:
 - Plugin: *hadesarchitect-cassandra-datasource*
 - Connection: *cassandra:9042*
- Template Variable (*\$topic*):
 - Type: Static list of stock tickers (AAPL, GOOGL, TSLA,...)
 - Used to filter all panels in the dashboard



IV. Serving Layer

Fear & Greed Index

- Visualization: Gauge
- Purpose: Display the latest market sentiment from 'twitter.topic_sentiment_avg'

Momentum: Bulls vs Bears

- Visualization: Time Series (Line Chart).
- Purpose: Compare bullish vs bearish tweet volume over time from 'twitter.topic_sentiment_avg'

Market Composition

- Visualization: Time Series (Stacked Area Chart)
- Purpose: Show overall discussion volume, visualize sentiment distribution

Stock Price History

- Visualization: Time Series (Line Chart)
- Purpose: Track stock price movement over time, correlate price trends with sentiment changes

Live Tweets Stream

- Visualization: Table
- Purpose: Display real tweet content in near real-time

Conclusion

The system has successfully implemented all core functionalities and demonstrated stable real-time data stream processing capabilities.

- Real-time sentiment analysis and stock price tracking.
- Managed with Terraform to ensure efficient and reproducible deployments.
- Built using Kafka, Spark, Cassandra, and Kubernetes.
- Grafana dashboards provide deep market insights.

This is
the end of
representation
